

Sales Performance & Forecasting Analysis

Python & Data Science Internship Assessment

Project Objective

This project performs an end-to-end analysis of the Superstore retail sales dataset. The analysis covers data cleaning, exploratory data analysis, statistical analysis, and machine-learning-based sales forecasting.

Technologies Used

Python 3, Pandas, NumPy, Matplotlib, Seaborn, Scikit-learn and Jupyter Notebook.

Dataset

Superstore Sales Dataset — Kaggle.

1. Project Description

In this project, I analyzed the Superstore sales data to understand sales patterns and tried to predict monthly sales using machine learning. The workflow includes inspection and cleaning of the raw dataset, exploratory visualizations, statistical summaries, predictive modeling, model evaluation, and business recommendations.

2. Data Cleaning & Preparation

The dataset contains 9,800 rows and 18 columns. The data was inspected for missing values, duplicate records, and incorrect data types.

There were 11 missing values in the Postal Code column. These values were replaced with the label 'Unknown' because Postal Code is a location identifier and replacing it with a numerical average would not be meaningful.

The Order Date and Ship Date columns were converted from text to datetime format. After conversion, both date columns contained zero missing values. No duplicate rows were found in the dataset.

3. Statistical Summary

Statistic	Sales
Mean	230.77
Median	54.49
Standard Deviation	626.65
Minimum	0.444
Maximum	22,638.48

Interpretation: The mean sales value of 230.77 is substantially higher than the median of 54.49. This indicates that a relatively small number of high-value orders strongly influence the average sales value.

4. Exploratory Data Analysis




Figure 1 — Monthly Sales Trend




Figure 2 — Sales by Product Category




Figure 3 — Sales by Region




Figure 4 — Correlation Heatmap




Figure 5 — Top 10 Products by Sales




Figure 6 — Actual vs Predicted Monthly Sales

5. Key Insights

- **Technology is the strongest category:** Technology generated the highest total sales at 827,455.87, followed by Furniture at 728,658.58 and Office Supplies at 705,422.33.
- **Regional performance varies:** The West region recorded the highest sales at 710,219.68, while the South region recorded the lowest at 389,151.46.
- **Sales fluctuate substantially over time:** Monthly sales ranged from 4,519.89 in February 2015 to 117,938.15 in November 2018.
- **Top product:** Canon imageCLASS 2200 Advanced Copier was the highest-selling product, generating 61,599.82 in total sales.

6. Predictive Modeling

A Linear Regression model was first developed using monthly time sequence information. The model achieved an RMSE of 23,619.29 and an R^2 score of 0.0332, indicating that a simple linear trend was insufficient to capture the variation in monthly sales.

An improved Random Forest Regression model was then developed using time-based and seasonal features including month number, year, month, and cyclical month features. The improved model achieved an RMSE of **16,094.10** and an R^2 score of **0.5511**.

Model	RMSE	R^2 Score
Linear Regression	23,619.29	0.0332
Random Forest	16,094.10	0.5511

7. Business Recommendations

1. Focus inventory and promotional activity on Technology products because this category generated the highest total sales.
2. Investigate the lower-performing South region to identify opportunities related to customer demand, product availability, pricing, and marketing.
3. Prioritize high-performing products such as the Canon imageCLASS 2200 Advanced Copier in inventory and promotional planning.
4. Prepare inventory, staffing, and marketing campaigns for high-demand periods based on historical monthly sales patterns.
5. Continue improving the forecasting model by incorporating additional variables such as customer segment, discount, shipping mode, product category, and lagged sales.

8. Conclusion

The Sales Performance & Forecasting Analysis identified important patterns in the Superstore dataset. Technology was the highest-performing category, while the West was the strongest region. Sales showed substantial variation across months, with November 2018 recording the highest monthly sales.

The Random Forest model performed substantially better than the initial Linear Regression model. The final Random Forest model achieved an R^2 score of 0.5511 and an RMSE of 16,094.10. The Random Forest model performed better than Linear Regression. Adding time and seasonal information helped the model predict sales better.

The model could be improved in the future by adding more useful features such as discounts, customer segments and previous months' sales.

9. GitHub Repository

GitHub Repository: **To be added after repository creation.**